

EXECUTIVE BRIEF

ProposalPilot

Technical Review · Database Audit · System Test · Competitive Analysis

Prepared by Perplexity Computer · May 3, 2026 · Confidential

Repository: github.com/Ryan-Richmond/RFPPProject · Database: Supabase ProposalPilot · Stack: Next.js 16 · Supabase · Perplexity AI

SECTION 1

Executive Summary

ProposalPilot is a well-architected, Perplexity-native AI platform designed to automate end-to-end government proposal response for small and mid-sized contractors. The core product thesis is strategically sound and differentiated: it targets a structurally underserved market — 500,000+ registered small businesses on SAM.gov priced out of every existing enterprise tool — with a full-pipeline workflow that no affordable competitor currently delivers.

The codebase is production-quality in architecture. The Supabase database is well-designed with comprehensive RLS, pgvector embeddings, and 19 carefully sequenced migrations, though several security advisories and performance gaps need attention before customer-facing scale. There is one critical deployment finding: the live URL at proposal-pilot.vercel.app hosts a different, unrelated application — the ProposalPilot GovCon codebase has not yet been deployed to Vercel.

Competitively, ProposalPilot is entering the market at the right moment. GovDash just raised \$30M in January 2026 at 16x revenue growth — confirming the category is real and expanding. Yet every funded competitor is enterprise-priced (\$30K+/year) or lacks the full-stack workflow. ProposalPilot's Perplexity-native architecture, success-fee pricing, and SMB focus represent a genuine structural gap in the market.

DIMENSION	ASSESSMENT	RATING
Product Vision & Thesis	GovCon-first, evidence-grounded, evidence-first — well-scoped for MVP	STRONG
Code Architecture	Next.js 16, Perplexity-native, clean service layer separation	STRONG
Database Design	Well-structured; security and performance advisories present	GOOD / FIX
Live Deployment	GovCon build not deployed; wrong app at proposal-pilot.vercel.app	CRITICAL GAP
AI Architecture	Perplexity Agent + Sonar + Embeddings; note: file imports say "gemini"	GOOD / FIX
Competitive Position	Strong SMB gap, unique success-fee model, right timing	STRONG
Sprint Readiness	19 migrations, 6 sprints planned, mock mode operational	ON TRACK

SECTION 2

Architecture & Code Review

2.1 Overall Architecture

ProposalPilot is built as a Next.js 16 App Router application with a Supabase backend, deployed via Vercel. The architecture cleanly separates concerns: UI pages, API routes, service functions, and AI calls are each in distinct layers. This is a sound, production-grade monorepo structure.

LAYER	TECHNOLOGY	NOTES
Frontend	Next.js 16, React 19, Tailwind CSS 4, shadcn/ui	Current, well-chosen stack
Backend	Next.js API routes (serverless)	Lean — no separate backend process
Database	Supabase (PostgreSQL 17.6 + pgvector + Storage + Auth)	Solid choice for this stage
AI Provider	Perplexity Agent API · Sonar API · Embeddings API	Single PERPLEXITY_API_KEY — clean
Hosting	Vercel (Next.js native)	Not yet deployed (see Section 4)
Embeddings	pplx-embed-context-v1-0.6b (1024 dimensions)	Migration 001 says 768d — mismatch to fix
Document Parsing	mammoth (DOCX), pdf-parse (PDF)	Handles core file types
Export	docx (npm package)	Word-compatible output

2.2 AI Layer Design

The CLAUDE.md architecture document specifies Perplexity as the single AI provider, with intelligent model routing: Claude Sonnet 4.6 for analysis/compliance, Claude Opus 4.6 for drafting, and Sonar-pro for web-grounded opportunity discovery. This is a compelling architectural decision — it enables real-time agency intelligence and live regulatory citations that no static content library can match.

Issue Found: The service files (rfp-analyzer, knowledge-base, compliance-checker) import AI functions from `@lib/ai/gemini` rather than `@lib/ai/perplexity`. The most recent commit confirms a Gemini API integration was merged on May 2, 2026. The `lib/ai/perplexity.ts` file contains the production-ready Perplexity implementation. Confirm which provider is currently active in production and align the import paths before deployment.

2.3 Agent Operation Logging

Every AI call writes a record to the `agent_operations` table (`status`, `model_used`, `input_summary`, `duration_ms`, `citations_count`). This is excellent practice — it enables observability, cost attribution by workspace, and debugging without re-running expensive calls. The pattern is consistent across all service files.

2.4 Mock Mode

The `AI_MODE=mock` environment variable disables all Perplexity API calls and returns deterministic fixture responses. This is the correct approach for development, demos, and CI. The mock data layer covers analysis, drafting, compliance, discovery, scoring, and embeddings. The Workspace page's Load Demo button seeds a complete end-to-end dataset.

2.5 Service Architecture Quality

The five core services (`rfp-analyzer`, `knowledge-base`, `proposal-drafter`, `compliance-checker`, `opportunity-scoring`) each follow a consistent pattern: fetch data from Supabase, call the AI layer, parse the structured JSON response with a fallback handler, store results, and return typed output. Error handling is present throughout. The knowledge base chunker uses a 500-word chunk with 50-word overlap — appropriate for the use case.

2.6 Findings & Recommendations

CRITICAL	AI Import Mismatch: Service files import from <code>lib/ai/gemini</code> but the <code>CLAUDE.md</code> spec says Perplexity-only. Resolve provider conflict and align all imports before production deployment.
CRITICAL	Embedding Dimension Mismatch: Migration 001 creates the <code>evidence_chunks.embedding</code> column as <code>vector(768)</code> but the Perplexity embeddings client targets 1024 dimensions. This will cause silent failures in vector insert operations. Run a migration to alter the column to <code>vector(1024)</code> .
IMPORTANT	match_evidence_chunks RPC: The <code>pgvector</code> similarity search calls <code>supabase.rpc("match_evidence_chunks")</code> but there is no migration creating this stored procedure. The fallback is a basic text search — ensure the RPC is created or the vector search will never run.
IMPORTANT	Chunking Strategy: Word-count chunking (500 words) does not respect document structure. Consider sentence-boundary or paragraph-boundary chunking to keep semantically coherent units together, which improves embedding quality.
SUGGESTION	Rate Limiting on AI Routes: No rate limiting is implemented on the <code>/api/opportunities/score</code> , <code>/api/proposals/[id]/draft</code> , or <code>/api/knowledge-base/index</code> routes. Add per-workspace per-day limits before multi-tenant deployment to control AI cost.
SUGGESTION	Export Quality: The docx export exists but citation metadata handling is not fully implemented in the proposal drafter. Ensure section-level citation IDs round-trip correctly into the exported document.

SECTION 3

Database Review (Supabase — ProposalPilot)

3.1 Schema Overview

The database has 19 migrations applied to a PostgreSQL 17.6 instance in us-west-2, currently ACTIVE_HEALTHY. The schema is comprehensive and purpose-built: it covers workspaces, tenant isolation, source document storage, knowledge base chunks with pgvector embeddings, solicitations, extracted requirements, compliance matrices, proposal drafts/sections/revisions, citations, outcomes, audit logs, SAM.gov opportunity ingestion, AI-powered scoring, pilot program management, and pipeline observability.

ENTITY GROUP	TABLES	ROW COUNT	RLS
Workspace & Auth	workspaces, workspace_members, workspace_invites	1 workspace, 1 member	✓ All
Company & Documents	company_profiles, source_documents, evidence_chunks	7 evidence chunks	✓ All
Proposals	solicitations, extracted_requirements, proposal_drafts, proposal_sections, proposal_section_revisions, citations, compliance_findings, compliance_matrix_entries	1 solicitation, 4 revisions	✓ All
Opportunity Discovery	opportunities, sam_opportunities, sam_opportunity_scores, sam_opportunity_score_history, sam_opportunity_requirement_profiles, sam_opportunity_recommendation_overrides	0 rows (not seeded)	✓ All
Scoring & Learning	opportunity_scores, scoring_model_configs, scoring_calibration_runs, score_distribution_snapshots, opportunity_feedback_events	0 rows	✓ All
Observability	agent_operations, pipeline_run_metrics, audit_logs, discovery_runs, opportunity_pipeline_dead_letters	1 pipeline run	✓ All
Pilot Program	pilot_workspaces, pilot_weekly_reports	0 rows	✓ All

3.2 What the Database Does Well

- Row Level Security (RLS) is enabled on every single table — complete tenant isolation from day one.
- The pgvector IVFFlat index on evidence_chunks.embedding with 100 lists is correctly configured for cosine similarity search.
- Dead-letter queue table (opportunity_pipeline_dead_letters) for failed pipeline events — production-grade resilience.
- Score history tracking (sam_opportunity_score_history) enables score drift analysis over time.

- Workspace invite system has expiry, revocation, email binding, and redemption tracking — well-designed.
- The sam_opportunities table has 9 dedicated indexes including GIN on the naics_codes array — well-optimized for filter-heavy queries.
- 19 sequential, well-named migrations with consistent DDL style.
- proposal_section_revisions includes actor_type (ai/user/system) and change_type — complete audit trail for AI edits.

3.3 Security Advisories

Supabase's security linter flagged two categories of warnings that should be resolved before customer-facing deployment:

Mutable Search Path on 9 Functions

The functions is_workspace_admin, handle_updated_at, handle_new_workspace, validate_workspace_invite, storage_object_workspace_id, handle_new_user_workspace_invite, redeem_workspace_invite, match_evidence_chunks, and is_workspace_member all have mutable search_path parameters. This exposes a schema injection risk where a malicious search_path could redirect function calls. Fix: add SET search_path = public, extensions to each function definition.

SECURITY CRITICAL: handle_new_user_workspace_invite(), handle_new_workspace(), is_workspace_admin(uuid), is_workspace_member(uuid), redeem_workspace_invite(text), and validate_workspace_invite(text, text) are all callable by the anon role as SECURITY DEFINER functions via the REST API. This means an unauthenticated user can call these functions. Remediation: REVOKE EXECUTE from anon on each function, or switch to SECURITY INVOKER. See: supabase.com/docs/guides/database/database-linter?lint=0028

Leaked Password Protection Disabled

Supabase Auth is not checking passwords against HaveIBeenPwned.org. Enable this in the Auth settings dashboard — it's a one-click change that prevents use of compromised credentials.

3.4 Performance Advisories

The performance linter identified 25+ unindexed foreign keys and 3 RLS initialization plan inefficiencies. These are low-impact at current data volumes (near zero rows) but will degrade as data scales.

ISSUE	AFFECTED TABLES	FIX
Unindexed FKs (25+)	audit_logs, citations, compliance_findings, compliance_matrix_entries, proposal_outcomes, proposal_section_revisions, workspace_invites, workspaces, and more	Add covering indexes on all FK columns listed in the advisor output
RLS re-evaluation per row	workspaces (workspace_owner_all), workspace_members (workspace_members_insert), sam_opportunities (sam_opportunities_write)	Replace auth.uid() with (select auth.uid()) in policy definitions

Unused Indexes (5)	idx_evidence_chunks_category, idx_proposal_drafts_solicitation, idx_audit_logs_workspace, idx_audit_logs_created, idx_opportunities_status	Monitor under real traffic; remove if still unused after 30 days of use
--------------------	--	---

3.5 Additional Schema Recommendations

- Add a composite index on (workspace_id, status) for solicitations and proposal_drafts — these are the most frequent dashboard query patterns.
- Consider partitioning agent_operations by created_at month once you have >100K rows — it will grow quickly under real usage.
- The evidence_chunks table stores naics_codes as text[] — add a GIN index on this array (same as the sam_opportunities table already has) to support fast NAICS-based capability matching.
- Add NOT NULL constraints on analysis_result in solicitations once the analyzer is stable — NULL rows can cause silent failures in the drafting workflow.
- workspace_members has member_email as a nullable column — consider enforcing NOT NULL with a trigger or application-level check to ensure invite redemption always captures the email.

SECTION 4

Live System Test

Critical Finding: The URL `proposal-pilot.vercel.app` hosts a different, simpler application — a freelance/B2B proposal generator with no GovCon features. It is completely unrelated to the ProposalPilot GovCon codebase in the RFPProject repository. The Vercel deployment for the RFPProject codebase was not found. The Vercel CLI confirmed "No deployments found" under the `rm14rich-8523` account.

4.1 The Wrong App — What Was Found

The app found at `proposal-pilot.vercel.app` is a narrow B2B proposal generator (for freelancers and consultants, not government contractors). It was tested thoroughly and the following was observed:

FEATURE	STATUS	NOTES
Landing page	WORKS	Clean design, fast load, EN/ES toggle
Signup / Login	WORKS*	Functional but email confirm link points to <code>localhost:3000</code> — BROKEN in production
Proposal creation flow	WORKS	15-20 sec generation, good AI quality output
Templates browser	WORKS	10+ templates, previews functional
PDF export / copy text	WORKS	Follow-up email generation included
GovCon features	NONE	No opportunities, knowledge base, compliance, workspace, team, drafting, or Computer Ops pages
Free tier limit	WARNING	Only 3 free AI generations — very low for trial conversion

4.2 Deployment Action Plan

To get the actual ProposalPilot GovCon build live, the following steps are required:

- 1 Resolve AI provider conflict (Gemini vs. Perplexity) in service imports before deploying.
- 2 Run the embedding dimension migration: `ALTER TABLE evidence_chunks ALTER COLUMN embedding TYPE vector(1024).`
- 3 Create a new Vercel project linked to the RFPProject repo, setting the root directory to `proposal-pilot/`.
- 4 Set all required environment variables in Vercel: `NEXT_PUBLIC_SUPABASE_URL`, `NEXT_PUBLIC_SUPABASE_PUBLISHABLE_KEY`, `AI_MODE`, `PERPLEXITY_API_KEY`.

- 5 Update the Supabase Auth email confirmation redirect URL from localhost:3000 to the production Vercel domain.
- 6 Apply all security fixes from Section 3 before opening to external users.
- 7 Run the Load Demo seed in mock mode to verify the full end-to-end workflow on the live URL.

SECTION 5

Competitive Analysis

5.1 Market Overview

The government contracting proposal market is undergoing rapid transformation driven by AI. A \$750B+ annual federal procurement market creates continuous demand for proposal automation, but the sector has historically been underserved relative to its complexity. The competitive landscape breaks into four strategic tiers:

Tier 1: Legacy Commercial RFP Tools	Responsive/RFPIO, Loopio, Qvidian	Built for commercial sales teams; no federal-specific compliance workflows; \$500–\$1,000+/user/month; enterprise-only
Tier 2: Established GovCon Platforms	Deltek GovWin IQ, Unanet ProposalAI, Privia/Xait	Purpose-built for federal contractors; strong BD workflows; \$12K–\$190K/year; not accessible to SMBs
Tier 3: Consulting-First	Shipley Associates	Methodology standard for proposals ("Shipley process"); no SaaS product; training/consulting only; no scalable tool
Tier 4: AI-Native GovCon	GovDash (\$30M raised), AutogenAI, GovEagle (YC), Rohirrim (\$15M)	Purpose-built AI; fastest growing; most still enterprise-priced (\$30K+/year) or lack end-to-end workflow

5.2 Key Competitor Profiles

GovDash		CLOSEST COMPETITOR
Funding	\$30M Series B — January 2026	
Pricing	Custom enterprise pricing — no SMB tier	
Strengths	End-to-end AI platform; 16x revenue growth YoY; ~200 customers; captures capture + proposal + BD in one platform	
Weaknesses	Enterprise-only; no success-fee model; no Perplexity-native live search intelligence	
AI	Proprietary AI fine-tuned on GovCon data	

Deltek GovWin IQ + Proposals		DOMINANT INCUMBENT
Funding	Roper Technologies (\$20B market cap)	
Pricing	\$12K/year entry — up to \$190K enterprise	
Strengths	6,000+ customers; most comprehensive opportunity intelligence; AI proposals launched 2025; deep integrations	
Weaknesses	AI proposals are new and immature; complex legacy system; no SMB pricing; no success-fee model	
AI	AI proposals added in 2025 as bolt-on; not native	

Unanet ProposalAI		DIRECT MID-MARKET COMPETITOR
Funding	Private equity (Riverside Partners)	
Pricing	~\$33/user/month; ProposalAI is add-on pricing	
Strengths	8 consecutive Inc. 5000; CRM + ERP + proposals in one platform; private cloud deployment for sensitive data	
Weaknesses	Seat-based pricing adds up fast; ProposalAI launched late 2024 — early stage; no SAM.gov discovery integration	
AI	ProposalAI launched late 2024; AI drafting in early maturity	

AutogenAI		ENTERPRISE-ONLY AI
Funding	\$22.3M raised	
Pricing	\$30K+/year (enterprise-only)	
Strengths	250+ enterprise customers; FedRAMP High authorized; UK-origin with strong public sector pedigree	
Weaknesses	Structurally inaccessible to SMBs; no SAM.gov opportunity discovery; no success-fee option	
AI	Enterprise-grade AI drafting with compliance focus	

Govly		COMPLEMENTARY
Funding	\$13.26M Series A	
Pricing	Free tier + \$3K/year Basic + Enterprise custom	
Strengths	Best GWAC/IDIQ opportunity discovery; network-effect moat; CMMC Level 2; free tier for SMBs	
Weaknesses	No proposal writing at all — opportunity discovery only; not a competitor, a potential integration partner	
AI	AI research tools for opportunity analysis, not generative drafting	

GovEagle		STARTUP COMPETITOR
Funding	\$2.5M (Y Combinator)	
Pricing	Not publicly disclosed	
Strengths	Shipley-aligned methodology; Microsoft Office native; YC network distribution	
Weaknesses	Proposal-stage only — no opportunity discovery; limited to Office ecosystem; narrow scope	
AI	AI drafting with Shipley compliance scoring	

5.3 ProposalPilot's Six Structural Differentiators

1 Perplexity-Native AI Architecture

Real-time web grounding via Sonar enables agency-aware proposals with live regulatory citations, budget intelligence, and incumbent research that no static content library can match. This is a capability no competitor currently offers.

2 End-to-End Workflow Coverage

SAM.gov discovery → parse → extract → draft → review → compliance check → export. Only GovDash comes close, and they are enterprise-only. No affordable platform delivers this full chain.

3 Success-Fee Pricing Model

The only platform in the market aligning incentives with contractor win rates. Eliminates \$20K–\$60K upfront cost barriers entirely. Directly addresses the #1 adoption objection for small contractors.

4 SMB Market Focus

500,000+ small businesses are registered on SAM.gov. Every existing AI-native competitor is structured for enterprises. ProposalPilot is purpose-built for the firms that need it most.

5 SAM.gov Native Integration

Built-in SAM.gov opportunity ingestion with AI scoring, NAICS matching, set-aside eligibility scoring, and capability gap analysis. Only ~3 platforms connect opportunity discovery + proposal writing — and none at SMB prices.

6 Perplexity Computer as Business Operator

Uniquely, Perplexity Computer is embedded as both the AI engine inside the product AND the operator running ProposalPilot's business operations (pipeline monitoring, customer support triage, competitive intelligence). This is a compounding advantage no competitor can replicate.

SECTION 6

Consolidated Recommendations

6.1 Critical (Fix Before Any External User Access)

CRITICAL	Resolve AI Provider Conflict The most recent commit integrated Gemini API, but CLAUDE.md specifies Perplexity-only. Service files import from lib/ai/gemini. Decide on one provider, remove the other, and align all imports. This is a fundamental architectural decision that must be resolved first.
CRITICAL	Fix Embedding Dimension Mismatch Migration 001 creates vector(768) but the Perplexity embeddings model produces 1024-dimensional vectors. This will cause silent failures on every knowledge base upload. ALTER TABLE evidence_chunks ALTER COLUMN embedding TYPE vector(1024) and rebuild the IVFFlat index.
CRITICAL	Create match_evidence_chunks RPC The knowledge base search calls supabase.rpc("match_evidence_chunks") but this stored procedure does not exist in any migration. Without it, every similarity search silently falls back to basic text search. Write and apply a migration creating this pgvector RPC.
CRITICAL	Fix SECURITY DEFINER Function Exposure Six SECURITY DEFINER functions are callable by the anon role via the REST API. REVOKE EXECUTE from anon on: handle_new_user_workspace_invite, handle_new_workspace, is_workspace_admin, is_workspace_member, redeem_workspace_invite, validate_workspace_invite.
CRITICAL	Deploy the Correct Codebase A different app is currently live at proposal-pilot.vercel.app. Create a new Vercel project pointing to the RFPPProject repo (root directory: proposal-pilot/). Set all required environment variables. Update Supabase auth redirect URLs.
CRITICAL	Fix Email Confirmation URL Whether in the current app or the GovCon build, the Supabase Auth email confirmation redirect is pointing to localhost:3000. Update the Site URL in Supabase Auth settings to the production domain.

6.2 Important (Before Pilot Program Launch)

IMPORTANT	Add Composite DB Indexes Add (workspace_id, status) indexes on solicitations and proposal_drafts. Add a GIN index on evidence_chunks.naics_codes. Fix the RLS re-evaluation issue on 3 tables by using (select auth.uid()) syntax.
IMPORTANT	Fix Mutable Search Path on 9 Functions Add SET search_path = public, extensions to is_workspace_admin, handle_updated_at, handle_new_workspace, validate_workspace_invite, storage_object_workspace_id, handle_new_user_workspace_invite, redeem_workspace_invite, match_evidence_chunks, is_workspace_member.
IMPORTANT	Enable Leaked Password Protection One-click toggle in Supabase Auth dashboard. Prevents use of compromised credentials — important before external users register.
IMPORTANT	Add Per-Workspace AI Rate Limits Add per-workspace daily limits on the /api/opportunities/score, /api/proposals/[id]/draft, and /api/knowledge-base/index routes. AI calls are expensive and unbounded rate could lead to unexpected costs with multiple concurrent workspaces.
IMPORTANT	Improve Chunking Strategy Switch from word-count chunking to sentence-boundary or paragraph-boundary chunking. This keeps semantically coherent units together and materially improves embedding quality and retrieval precision.
IMPORTANT	Harden Citation Round-Trip Verify that citation IDs from proposal sections persist correctly through to the Word export. Evidence-traceability is a core product promise — it must survive export.

6.3 Strategic Suggestions

STRATEGY	Position Against GovDash Directly GovDash's \$30M Series B validates the AI GovCon proposal market. Their enterprise pricing creates a clean SMB wedge for ProposalPilot. Lead messaging should explicitly position around the pricing gap: "Enterprise AI capabilities at small business prices."
----------	---

STRATEGY	Govly Partnership Opportunity Govly does opportunity discovery; ProposalPilot does proposal response. These are adjacent, non-competing workflows targeting the same customer. An integration partnership or co-marketing relationship would compound reach for both.
STRATEGY	Lean into the Shipley Framework Shipley Associates has defined the methodology standard for government proposals but has no scalable SaaS product. ProposalPilot's structure (requirements extraction, compliance matrix, review workflow) is highly Shipley-compatible. Marketing around Shipley methodology alignment could accelerate trust with experienced proposal writers.
STRATEGY	Federal IT Services First The PRD's suggested beachhead of IT services federal contractors is correct. NAICS codes 541512, 541511, and 541519 represent the largest volume of competitive federal IT procurements. Seed the SAM.gov opportunity discovery with these NAICS codes for the pilot cohort.
STRATEGY	Pilot Program: 5 Workspaces, 90 Days The pilot_workspaces table and pilot program infrastructure in the database is production-ready. Recruit 5 small IT services contractors for a 90-day pilot. Track: time from upload to first draft, requirement extraction accuracy, compliance score before export, and proposals submitted. These metrics will be the foundation of the GTM story.
STRATEGY	Success Fee Legal Review The success-fee monetization model is strategically compelling but needs legal grounding. Define the cap, reporting mechanism, enforcement approach, and customer-friendly contract language before exposing it in the product. The PRD correctly notes this should not block the MVP but should be designed before the pilot.

CLOSING

Overall Assessment

ProposalPilot is one of the most well-considered GovCon AI products reviewed. The market thesis is sharp, the architecture is sound, the database is mature, and the competitive timing is right. The path to pilot-ready is not long — the critical blockers (AI provider conflict, embedding dimension mismatch, missing RPC, security advisories, and deployment) can each be resolved in a focused engineering sprint. The Perplexity-native architecture and success-fee pricing model are genuine moats. Fix the foundation, deploy the product, and run the pilot.

The federal procurement market represents trillions of dollars of annual spend. Small contractors — the ones ProposalPilot is designed for — submit hundreds of proposals each year at significant labor cost, with win rates that barely break 30%. Every percentage point of improvement in proposal quality, coverage, and compliance translates directly into revenue. That is a product that sells itself once demonstrated on a real bid.

Prepared by Perplexity Computer · May 3, 2026 · Sources: GitHub/Ryan-Richmond/RFPPProject, Supabase ProposalPilot (fzukcuiwiexrfquuzqii), Vercel CLI, Live browser test at proposal-pilot.vercel.app, Competitive research: Govly, GovDash, Deltek, Unanet, AutogenAI, GovEagle, Rohirrim, RFPIO/Responsive, Loopio, Qvidian, Shipley Associates, pWin.ai, McCarren AI